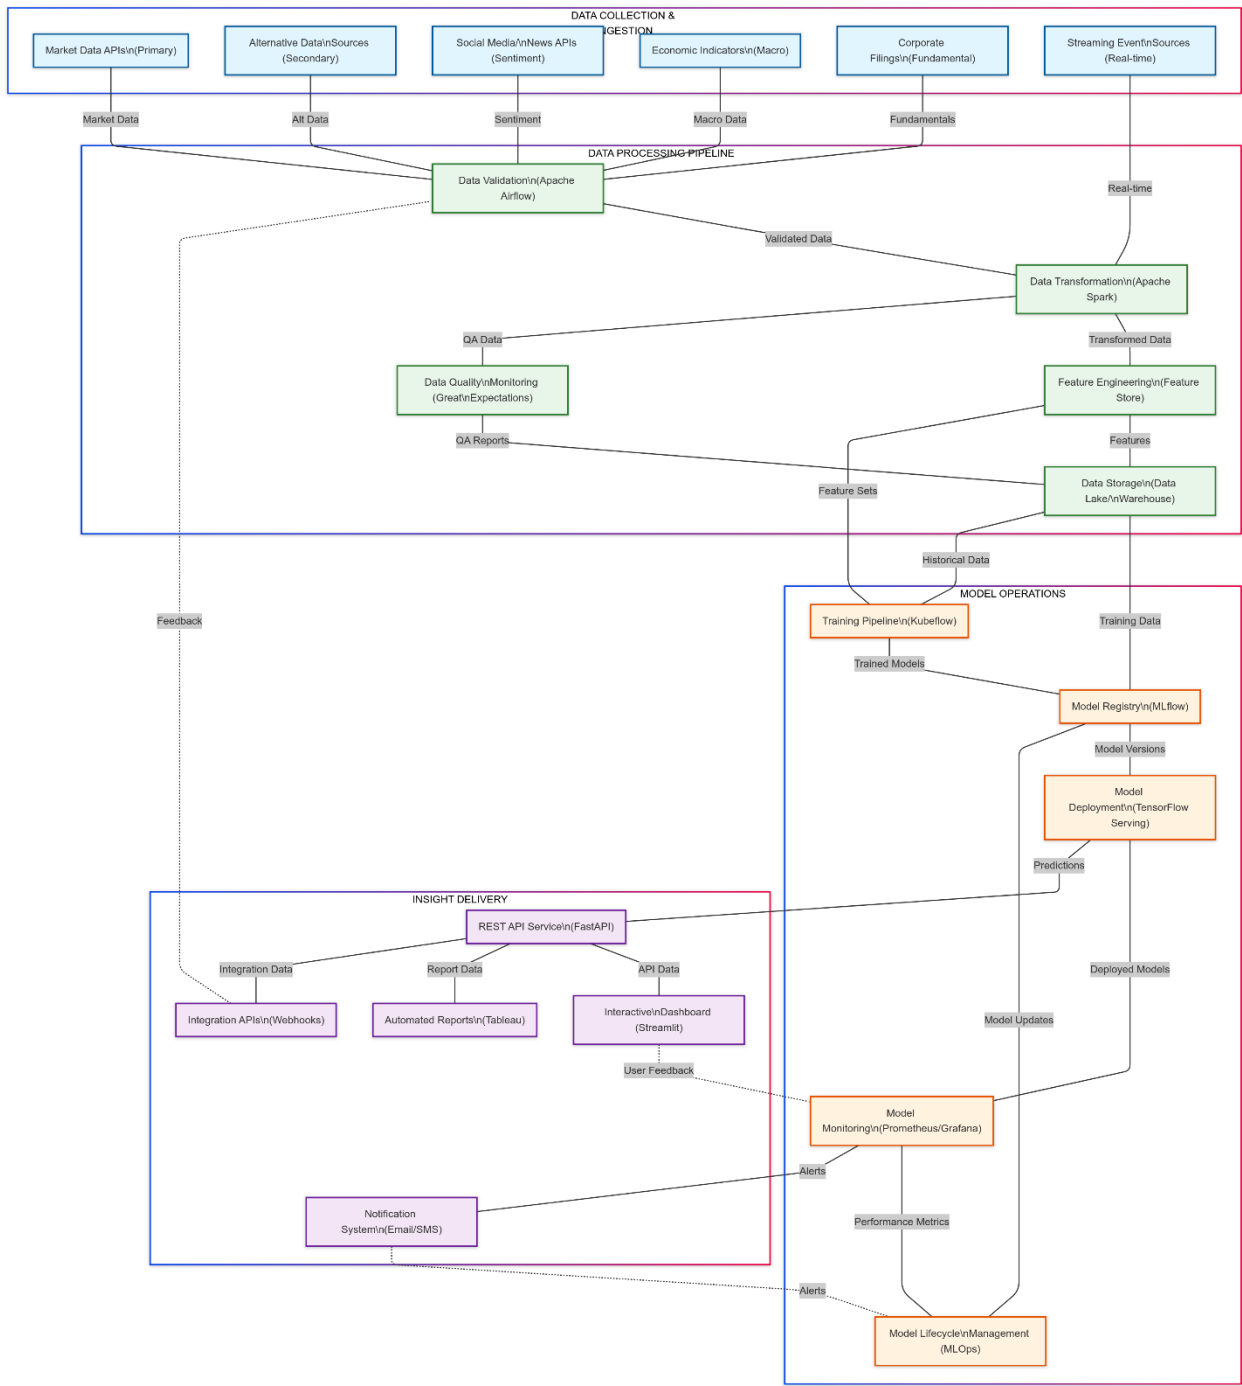


# End-to-End System Design for Stock Price Prediction

## 1. System Architecture Diagram



## 2. Component Justification

### Data Collection & Ingestion

#### *Market Data APIs (Primary)*

- **Technology:** API integrations with providers like Alpha Vantage, Yahoo Finance, or IEX Cloud
- **Justification:** These established providers offer reliable, structured market data with historical depth
- **Tradeoffs:** Subscription costs vs. data quality and reliability; rate limits may require batching strategies

#### *Alternative Data Sources (Secondary)*

- **Technology:** Specialized data providers (Quandl, Refinitiv, Bloomberg)
- **Justification:** Provides unique insights beyond traditional market data, creating potential prediction edges
- **Tradeoffs:** Higher cost; requires extensive validation and preprocessing

#### *Social Media/News APIs (Sentiment)*

- **Technology:** Twitter API, News APIs (NewsAPI, GDELT)
- **Justification:** Sentiment analysis can provide leading indicators of market movements
- **Tradeoffs:** Noisy data requiring sophisticated NLP processing; variable reliability

#### *Economic Indicators (Macro)*

- **Technology:** FRED API (Federal Reserve Economic Data)
- **Justification:** Macroeconomic conditions significantly influence market movements
- **Tradeoffs:** Lower update frequency; requires careful feature engineering for relevance

#### *Corporate Filings (Fundamental)*

- **Technology:** SEC EDGAR API
- **Justification:** Fundamental data provides long-term valuation context
- **Tradeoffs:** Infrequent updates; requires complex text processing for extraction

#### *Streaming Event Sources (Real-time)*

- **Technology:** WebSocket connections to exchange data feeds
- **Justification:** Enables real-time model updates and immediate trading signals
- **Tradeoffs:** Infrastructure complexity; higher bandwidth and processing requirements

### Data Processing Pipeline

#### *Data Validation (Apache Airflow)*

- **Technology:** Apache Airflow for workflow orchestration

- **Justification:** Reliable scheduling and monitoring of data pipelines with dependency management
- **Tradeoffs:** Requires infrastructure management; steeper learning curve than simpler schedulers

#### *Data Transformation (Apache Spark)*

- **Technology:** Apache Spark for distributed data processing
- **Justification:** Handles large-scale data processing efficiently with built-in ML capabilities
- **Tradeoffs:** Resource intensive; requires specialized knowledge for optimization

#### *Feature Engineering (Feature Store)*

- **Technology:** Feast or Tecton feature store
- **Justification:** Centralizes feature computation logic; ensures consistency between training and serving
- **Tradeoffs:** Additional system complexity; integration challenges with existing infrastructure

#### *Data Quality Monitoring (Great Expectations)*

- **Technology:** Great Expectations for data validation
- **Justification:** Ensures data meets quality expectations before entering the model pipeline
- **Tradeoffs:** Requires upfront investment in defining expectations; maintenance overhead

#### *Data Storage (Data Lake/Warehouse)*

- **Technology:** Hybrid solution with S3/Delta Lake (data lake) and Snowflake (data warehouse)
- **Justification:** Lake provides raw storage flexibility; warehouse enables efficient analytical queries
- **Tradeoffs:** Cost management challenges; requires data governance strategy

### **Model Operations**

#### *Model Registry (MLflow)*

- **Technology:** MLflow for model tracking and versioning
- **Justification:** Open-source solution with comprehensive tracking capabilities and broad framework support
- **Tradeoffs:** Requires additional infrastructure for scaling; UI limitations for complex use cases

#### *Training Pipeline (Kubeflow)*

- **Technology:** Kubeflow Pipelines for model training orchestration
- **Justification:** Containerized, reproducible ML workflows with Kubernetes scalability

- **Tradeoffs:** Complex setup; steep learning curve; requires Kubernetes expertise

#### *Model Deployment (TensorFlow Serving)*

- **Technology:** TensorFlow Serving for model serving (or similar based on model type)
- **Justification:** High-performance, production-grade serving with versioning support
- **Tradeoffs:** Framework-specific; requires adaptation for non-TensorFlow models

#### *Model Monitoring (Prometheus/Grafana)*

- **Technology:** Prometheus for metrics collection, Grafana for visualization
- **Justification:** Industry-standard monitoring stack with excellent alerting capabilities
- **Tradeoffs:** Requires instrumentation of services; additional infrastructure to maintain

#### *Model Lifecycle Management (MLOps)*

- **Technology:** Custom MLOps framework integrating above components
- **Justification:** Automates model retraining, promotion, and rollback based on performance metrics
- **Tradeoffs:** Development effort; organizational processes must adapt to automated lifecycle

### **Insight Delivery**

#### *REST API Service (FastAPI)*

- **Technology:** FastAPI for prediction endpoints
- **Justification:** High-performance Python framework with automatic documentation and validation
- **Tradeoffs:** Requires API management for authentication, rate limiting, etc.

#### *Interactive Dashboard (Streamlit)*

- **Technology:** Streamlit for analyst-facing dashboards
- **Justification:** Rapid development of interactive visualizations with Python-based implementation
- **Tradeoffs:** Performance limitations for very complex visualizations; customization constraints

#### *Automated Reports (Tableau)*

- **Technology:** Tableau for scheduled client reports
- **Justification:** Enterprise-grade visualizations with rich formatting and distribution capabilities
- **Tradeoffs:** License costs; separate system from the core ML pipeline

#### *Notification System (Email/SMS)*

- **Technology:** Amazon SES for email, Twilio for SMS

- **Justification:** Reliable, scalable alerting for critical predictions and model issues
- **Tradeoffs:** Additional integration points; regulatory compliance requirements for financial alerts

#### *Integration APIs (Webhooks)*

- **Technology:** Webhook architecture for third-party system integration
- **Justification:** Enables push-based integration with trading platforms and external systems
- **Tradeoffs:** Requires robust retry mechanisms; security challenges for financial data

### 3. Data Flow Explanation

#### Batch vs. Streaming Decisions

##### *Batch Processing (Daily)*

- **Implementation:** Scheduled Airflow DAGs run daily after market close
- **Data Scope:** Complete daily market data, fundamental updates, economic indicators
- **Justification:** Most financial data sources update on a daily cadence; enables comprehensive retraining
- **Key Processes:**
  - Full feature recalculation (technical indicators across multiple timeframes)
  - Model performance evaluation against daily outcomes
  - Periodic model retraining with extended historical data

##### *Near-Real-Time Processing (Minute-level)*

- **Implementation:** Spark Structured Streaming processes market data in micro-batches
- **Data Scope:** Price, volume updates throughout trading day
- **Justification:** Enables intraday prediction updates without full real-time complexity
- **Key Processes:**
  - Incremental feature updates (e.g., updating moving averages)
  - Model inference with updated features
  - Trading signal generation based on prediction thresholds

##### *Real-Time Processing (Event-based)*

- **Implementation:** Kafka streams process significant market events
- **Data Scope:** Breaking news, unusual market movements, trading halts
- **Justification:** Critical events require immediate system response regardless of scheduled updates
- **Key Processes:**
  - Event classification and prioritization
  - Targeted feature recalculation for affected securities
  - Alert generation for significant prediction changes

## Data Transformation Stages

### 1. Ingestion Stage

- Raw data collection from various sources
- Initial format standardization and timestamp normalization
- Preliminary data quality checks

### 2. Enrichment Stage

- Join data from multiple sources (market data + news sentiment + economic indicators)
- Compute derived metrics (volatility measures, liquidity indicators)
- Apply domain-specific transformations (adjustments for splits, dividends)

### 3. Feature Engineering Stage

- Calculate technical indicators as identified in the EDA (MA, RSI, MACD, Bollinger Bands)
- Extract temporal features (day-of-week, month, etc.)
- Generate market regime indicators based on historical patterns

### 4. Model-Ready Transformation Stage

- Scale/normalize features as required by model specification
- Create time-lagged features for sequence modeling
- Generate target variables at different prediction horizons (1-day, 5-day, 20-day)

### 5. Output Transformation Stage

- Convert model predictions to actionable insights
- Calculate confidence intervals and prediction risks
- Format data for visualization and reporting systems

## System Interaction Points

### 1. Data Provider Interfaces

- API authentication and rate limit management
- Scheduled and event-triggered data pulls
- Error handling and retry mechanisms

### 2. Data Processing to Model Operations

- Feature store serves as the primary handoff point
- Model training triggered by data quality verification
- Model registry receives newly trained candidates

### 3. Model Operations to Insight Delivery

- Model serving layer provides prediction endpoints to API service
- Monitoring systems feed performance metrics to dashboards

- Evaluation results determine model promotion decisions

#### 4. **Insight Delivery to End Users**

- REST API provides programmatic access for automated trading systems
- Interactive dashboards for analysts to explore predictions
- Notification system for time-sensitive alerts
- Scheduled reports for management and clients

#### 5. **Feedback Loops**

- Actual market outcomes flow back to evaluation systems
- User interactions with predictions captured for model improvement
- System performance metrics inform infrastructure scaling decisions

### 4. Challenge Analysis and Mitigation Approaches

#### Challenge 1: Data Quality and Consistency

**Problem:** Financial data often contains errors, gaps, or inconsistencies that can significantly impact model performance.

**Mitigation Approach:**

- Implement multi-source validation by cross-checking key metrics across different data providers
- Deploy automated anomaly detection specifically tuned for financial time series
- Create a data lineage system to track provenance of all features
- Establish clear data SLAs with notification thresholds for quality issues
- Develop graceful degradation strategies when primary data sources fail

#### Challenge 2: Model Drift in Changing Market Conditions

**Problem:** Financial markets exhibit regime changes that can rapidly invalidate previously effective models.

**Mitigation Approach:**

- Implement continuous model performance monitoring comparing offline metrics to online performance
- Develop market regime detection algorithms to automatically trigger model switching
- Maintain ensemble models optimized for different market conditions
- Create automated A/B testing framework for continuous model evaluation
- Establish clear thresholds for model retraining and retirement based on performance degradation patterns

### Challenge 3: Latency Requirements for Trading Applications

**Problem:** Trading decisions often require sub-second response times, challenging traditional ML pipelines.

**Mitigation Approach:**

- Optimizing critical path feature calculations for minimal computation.
- Implement feature caching with appropriate invalidation strategies.
- Deploy models with TensorRT or ONNX Runtime optimization for inference acceleration.
- Utilize edge deployment for prediction services physically close to trading infrastructure.
- Develop tiered service architecture with fast approximate predictions followed by refined analysis.

### Challenge 4: Regulatory Compliance and Explainability

**Problem:** Financial services face strict regulatory requirements around model transparency and audit capabilities.

**Mitigation Approach:**

- Implement comprehensive model documentation automated within the MLOps pipeline.
- Integrate SHAP or LIME explainability tools directly into the prediction service.
- Maintain immutable audit logs of all model versions, training data, and predictions.
- Develop standardized model cards that document limitations and appropriate use cases.
- Create a compliance-friendly feature importance visualization system for end users.

### Challenge 5: Scaling Economics and Resource Optimization

**Problem:** Financial data processing and modeling can become prohibitively expensive on a scale.

**Mitigation Approach:**

- Implement tiered data storage with hot/warm/cold zones based on access patterns.
- Develop automated resource scaling tied to market hours and volatility conditions.
- Create model complexity budgets based on prediction value vs. computing cost.
- Implement feature selection optimization in the training pipeline to reduce dimensionality.
- Develop cost attribution system to align infrastructure expenses with business value.